RecourseLLM

Give AI a World.

Not a Window.

LLMs are using context windows in the wrong way.
Recourse does it better.

Pre-Seed Investment Opportunity

Spring 2026

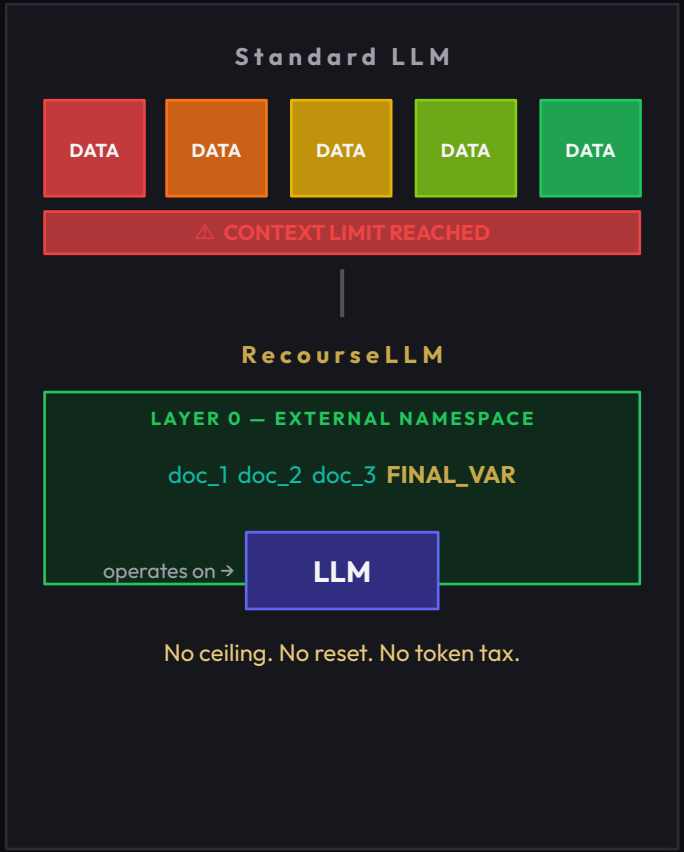

Every Enterprise AI Project Hits the Same Wall.

Not because the AI isn't smart enough. Because the problem is bigger than the window.

Documents too large to load

M&A data rooms. Legal portfolios. Financial filings. A 4,000-document corpus hits the ceiling immediately.

Sessions that forget everything

Every new conversation starts cold. No persistent state. No memory of yesterday's analysis.

Costs that scale with complexity

100K+ tokens per query. Multiply by 50,000 monthly operations. The bill becomes the bottleneck.

The Context Window Is a Control Plane, Not a Data Plane.

STANDARD LLM APPROACH		RECOURSELLM APPROACH
1 Document → loaded into context window		1 Document → REPL variable in Layer 0
2 Context window grows with data size		2 Context window stays lean — metadata only
3 Token cost scales with problem complexity	VS	3 Token cost independent of data scale
4 Session ends → all state lost		4 State persists across sessions indefinitely
5 Parallel tasks = redundant data loading		5 Parallel tasks share external namespace

RecourseLLM

Five Layers. One Principle.

Data never enters the model.
The model operates on data.

Completion_Signal
The universal completion signal. When the model writes the completion signal the task is done — regardless of output type.

sub_RLM_batch()
Parallel task delegation. Tasks fan out across sub-models sharing the same Layer 0 data — zero redundant transmission.

The Surgeon Analogy
A world-class surgeon doesn't memorize every chart, filing, and record before operating. She queries what she needs. Her working mind stays focused on judgment. RecourseLLM works the same way.

4 BT Orchestration
Reactive behavior trees. Parallel execution. Long-horizon task coordination.

3 Skill Registry
Composable skills that self-register at runtime. The system grows smarter through use.

2 Tool Registry
42+ primitives. Zero-token discovery. Results write to variables, never to context.

1 SafeREPL
LLM generates Python code. Code executes against Layer 0. Completion_Signal signals end.

0 External Working Memory
SemanticREPLDict + Vector DB. Every variable auto-indexed. LLM sees metadata only.

RecourseLLM

Token Reduction Is Not the Point.

It's the byproduct of doing this right.

TOKEN COST PER COMPLEX ENTERPRISE TASK

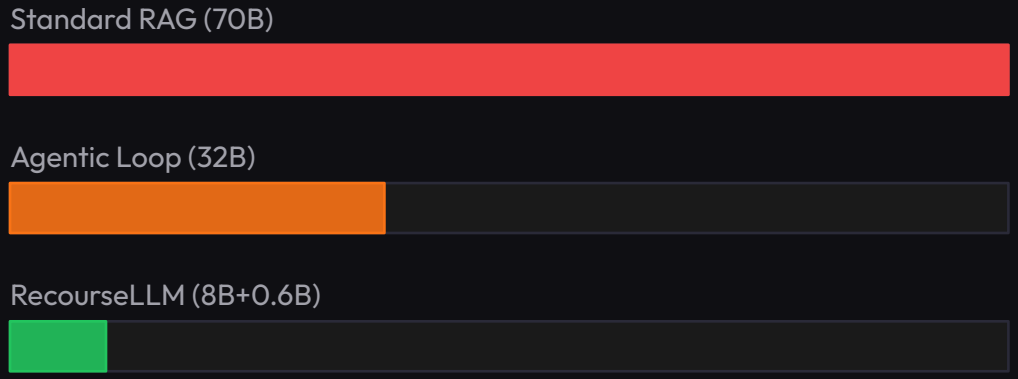

98% token reduction on complex tasks. At 1B+ monthly tokens, that's a projected $120M+ in annual infrastructure savings at scale.

98%
Token Reduction
vs Standard RAG

80%
Energy Reduction
at 1B+ monthly tokens

4.2×
Compute ROI
increase at enterprise scale

RecourseLLM

The Living World Problem. AI Has Never Solved It.

CURRENT STATE

✗ NPCs forget everything
Each conversation starts cold. No persistent memory.

✗ Rules hallucinate
LLMs generate plausible-sounding rule violations constantly.

✗ Worlds reset
Every session, the world's state is gone. Nothing carries forward.

✗ Scale is intractable
200 NPCs with individual state? Computationally impossible today.

WITH RECOURSELLM

✓ NPCs remember everything
Each NPC is a persistent Layer 0 namespace. Memory accumulates forever.

✓ Rules execute as code
D&D rulebook becomes deterministic REPL tools. No hallucination possible.

✓ Worlds persist indefinitely
World state lives in Layer 0. Sessions pick up exactly where they ended.

✓ Scale is native
sub_RLM_batch ticks 200 NPCs simultaneously. Cost is flat regardless.

"You could chat with an NPC in the morning — they wouldn't know the day's events yet. Chat with them that evening — they recall their entire day."

RecourseLLM

Gaming Leads. Enterprise Follows.

PHASE 1

NPC Co-Pilot

Months 1–6

- Deploy NPC living world demo
- Beta with Ultimate Game Master community (1M+ reach)
- D&D 5e 2024 full ruleset as executable skill library
- Prove: persistent NPCs, deterministic rules, temporal memory

PHASE 2

Game Studio Partnerships

Months 6–18

- Target mid-size RPG studios: Larian, Owlcat, ambitious indie
- Integration API for NPC behavior engine
- Proof: what Baldur's Gate 3 does in scripting, we do in real-time
- First licensing or integration revenue

PHASE 3

Enterprise Platform

Months 12–24

- Legal tech: M&A due diligence portfolios
- Financial services: longitudinal behavioral analysis
- Engineering: large codebase operations at scale
- Same architecture. Different world. Larger checks.

The Gaming community is the proof of concept. The game industry is the market. Enterprise is the platform.

RecourseLLM

The Markets.
One Architecture.

$3B+

TTRPG & Tabletop

D&D alone: $1B+. Growing 15% YoY. Passionate, tech-forward audience.

$200B+

Video Game AI

Global games industry. AI NPC behavior is the unsolved frontier for every major RPG studio.

$50B+

Enterprise AI Tooling

Legal, financial, engineering orgs hitting context ceiling walls today. Budgets are real.

Prove it in gaming. Own the architecture across verticals. Total addressable market exceeds $250B.

RecourseLLM

This Is Not a Feature Comparison. It's a Paradigm Comparison.

Capability	OpenAI / Anthropic	LangChain / RAG	RecourseLLM
Data scale ceiling	Hard context limit	Chunked (lossy)	**None ✓**
Session persistence	None	Manual / partial	**Native ✓**
Rule determinism	Probabilistic	Probabilistic	**Code-enforced ✓**
Parallel execution cost	Scales with data	Scales with data	**Data-independent ✓**
Self-improving skills	Not supported	Not supported	**Runtime registration ✓**
Token cost at scale	Linear growth	Linear growth	**Decoupled ✓**
Fine-tuning required	For domain knowledge	For domain knowledge	**Not needed ✓**

Anthropic and OpenAI are infrastructure providers. RecourseLLM is built on top of them — not competing with them.

RecourseLLM

Many Industries. One Shared Limit.

LEGAL — M&A DUE DILIGENCE

Problem: 4,000-document data rooms. Context window hits its ceiling on document 3.

Solution: Entire data room in Layer 0. Cross-portfolio queries run without loading. 'Find every clause conflicting with the master agreement' — answered completely.

Months of analyst time → hours

FINANCIAL — CLIENT INTELLIGENCE

Problem: 20 years of behavioral data. Standard AI can't hold enough context to find the patterns that matter.
Solution: Longitudinal analysis across full history. Which clients show spending shift signals? Ranked today, not approximated.

$120M+ at 1B token monthly scale

ENGINEERING — LEGACY CODEBASE

Problem: 800,000 lines of COBOL. No AI tool can hold the whole codebase in context. Refactoring is manual.
Solution: Entire codebase in Layer 0. sub_RLM_batch maps all dependencies simultaneously. Migration code generated module by module at scale.

18-month project → weeks

RecourseLLM

A Rare Combination in AI Infrastructure.

RecourseLLM Architect

Designed and built the five-layer architecture. Deep technical ownership of every component from Layer 0 external memory to Layer 4 behavior tree orchestration.

Gaming

Built one of the first Facebook games ever built (MouseHunt, 20 years old in 2027). Alpha partner to develop payment systems, distributed asset systems, and localization systems.

Military Constrained-Systems Architectures

Active in patenting, building and deploying military products alongside JTF2 (Elite Forces Engineering) and DOD (Canada's Department of Defense).

Enterprise Strategy Background

3 Previous exits in the gaming, music, and military defense/ robotics industries.

Technical depth. Community credibility. Enterprise relationships. Proven ability to raise.

RecourseLLM

Built. Running. Demonstrable.

COMPLETE

Architecture Document

Five-layer stack fully specified. Layer 0 external working memory, SafeREPL execution, tool registry, skill library, and BT orchestration all designed and documented.

FILING

Provisional Patent

Four novel mechanisms filed: SemanticREPLDict auto-indexing, variable-name task delegation, FINAL_VAR completion signal, BT orchestration with shared namespace.

ACTIVE

Community Distribution

The Ultimate Game Master: 4000 Beta NPC-builders (Game Masters), 'Brain' AI podcast audience. Direct access to beta audience for RPG product and technical credibility for enterprise.

LIVE

Website Live

recoursellm.com positioned and live. Architecture page, performance benchmarks, enterprise access flow. Rewrite in progress to reflect paradigm positioning.

IN PROGRESS

Enterprise Conversations

Initial outreach to legal tech and financial services prospects identified. Pilot structure defined: 90 days, real data, agreed success metrics, no cost.

RecourseLLM

Pre-Seed Round.

$1.5 Million

Pre-Seed · SAFE or Priced Round

USE OF FUNDS

40% **Engineering/Compute**
Core architecture hardening, API layer, SDK, developer experience

30% **GTM — Gaming**
NPC product launch, studio partnership outreach, community activation

20% **Enterprise Pilots**
3–5 pilot deployments: legal tech and financial services

10% **IP & Operations**
Legal, infrastructure

MILESTONES THIS ROUND ACHIEVES

M3 Public beta. 500+ registered beta.

M6 First game studio pilot agreement signed.

M9 3 enterprise pilots live. First paying customer.

M12 Product-market fit validated in 2 verticals.

M18 Skill library has 200+ registered skills.

RecourseLLM

Give AI a World. Not a Window.

The context window is the defining constraint of modern AI. Every enterprise, game studio, and developer building with LLMs is either working around it or being defeated by it. RecourseLLM is the first architecture that removes it from the equation entirely — not by making it bigger, but by making data scale irrelevant to intelligence.

Book a Demo · recoursellm.com